

Building data applications with Apache Druid

October 2020

Gian Merlino, Cofounder & CTO, Imply



VIRTUAL
DRUID SUMMIT

Who am I?

Gian Merlino



Committer & PMC chair at   druid


Cofounder at  (we're hiring!)

Building data apps



VIRTUAL
DRUID SUMMIT

Druid powers data apps



Druid in the wild

How Netflix uses Druid for Real-time Insights to Ensure a High-Quality Experience



Netflix Technology Blog

Follow

Mar 3 · 9 min read



By [Ben Sykes](#)

100+ billion rows/day

1+ trillion rows, 1+ year retained

100s of servers

sub-second to few seconds query latency

mix of streaming and batch ingest

Interactive Analytics at MoPub: Querying Terabytes of Data in Seconds

By [Aaron Rolett](#) and [Shravana Krishnamurthy](#)

Wednesday, 3 July 2019    

How Druid enables analytics at Airbnb

Realtime and batch analytics at Airbnb and the role Druid plays in our analytics system architecture



Pala

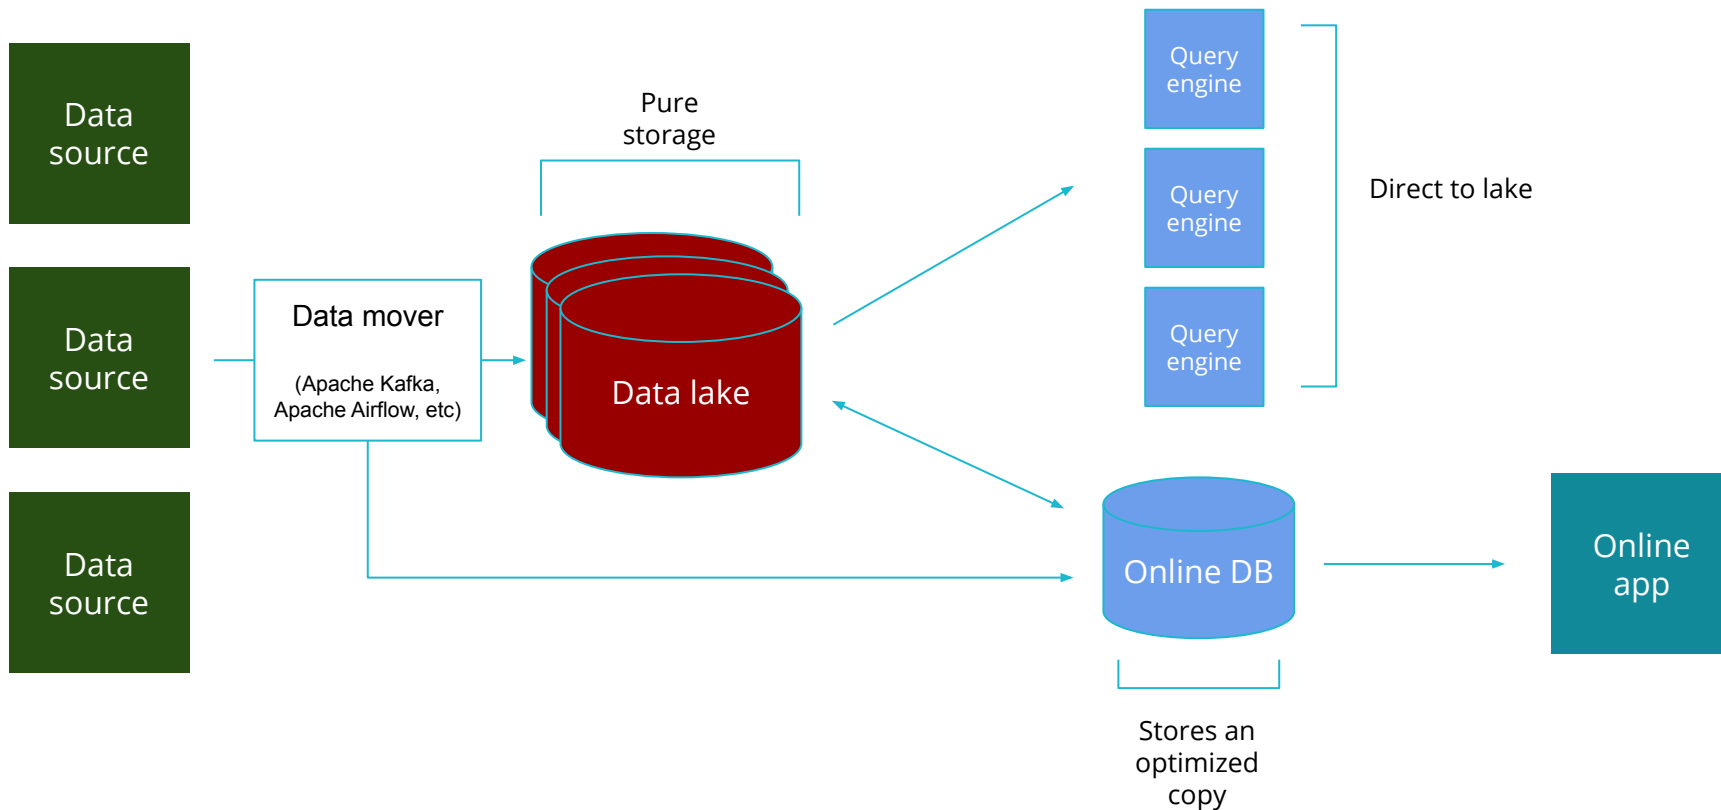
Follow

Nov 13, 2018 · 10 min read

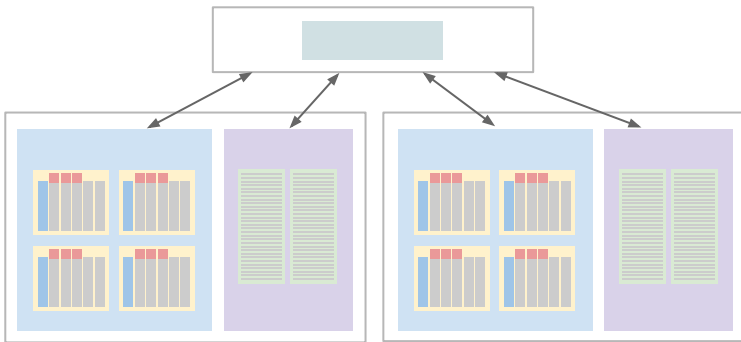


By [Pala Muthiah](#) and [Jinyang Li](#)

Online DB pattern



Druid as an online DB

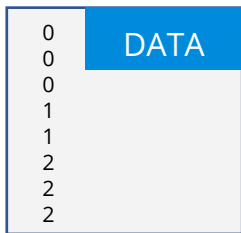


- Scale-out, fault-tolerant architecture
- No downtime for software updates
- No downtime for data management
- Heavily optimized storage format
- Integrated storage format and query engine

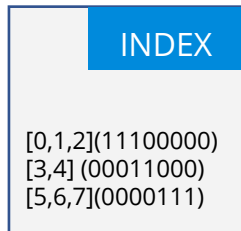
Data apps

- Interactive query speeds
- Always online
- Fresh data from streams
- Quality of service
- Price/performance

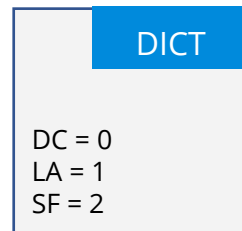
Interactive query speeds




Compression



Secondary indexes

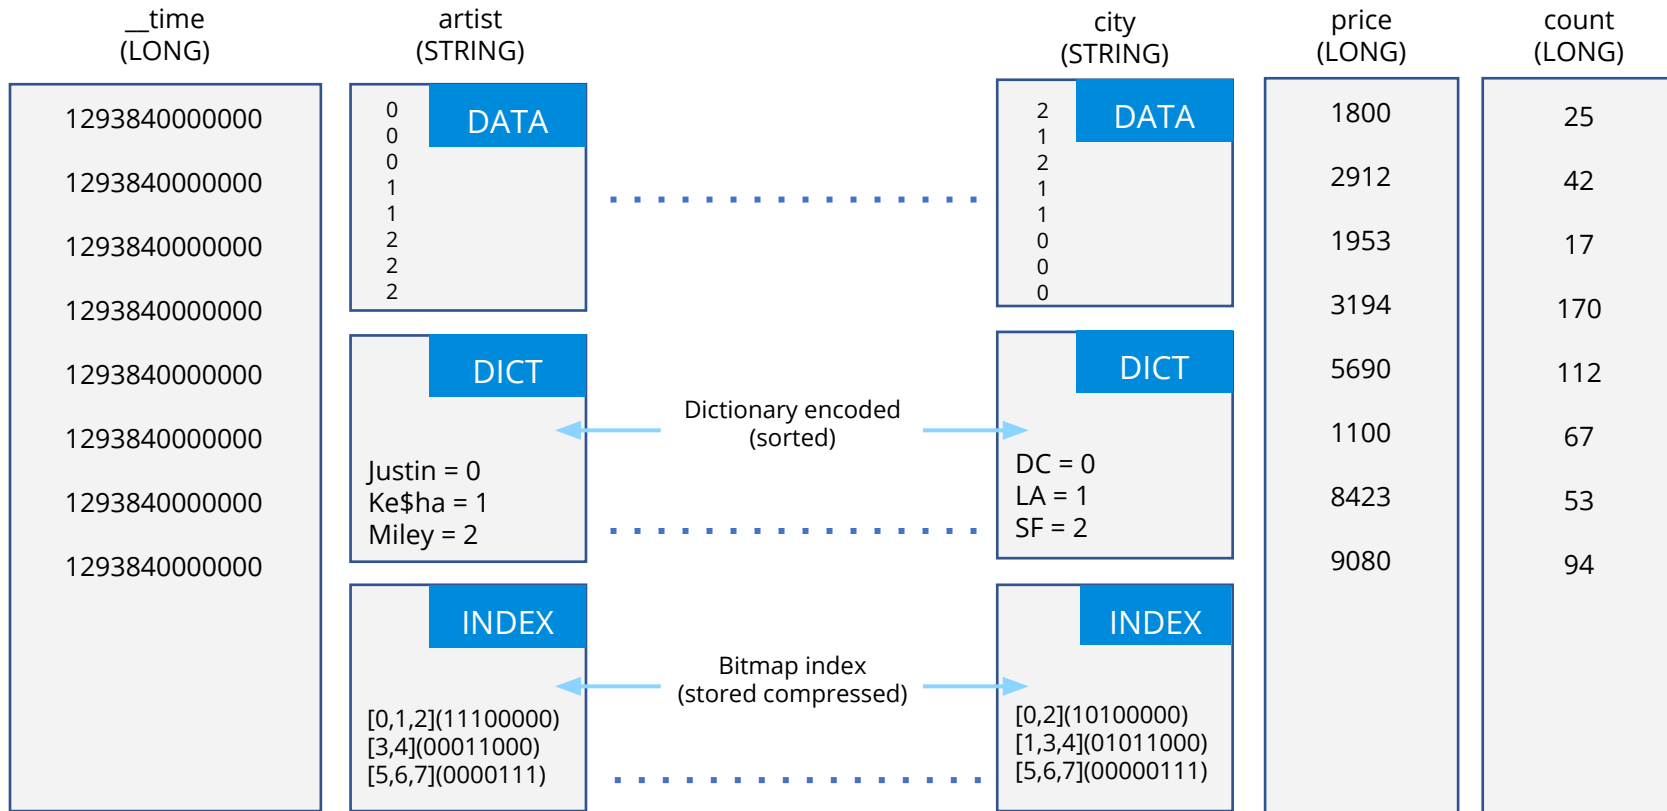


Operate on
compressed data

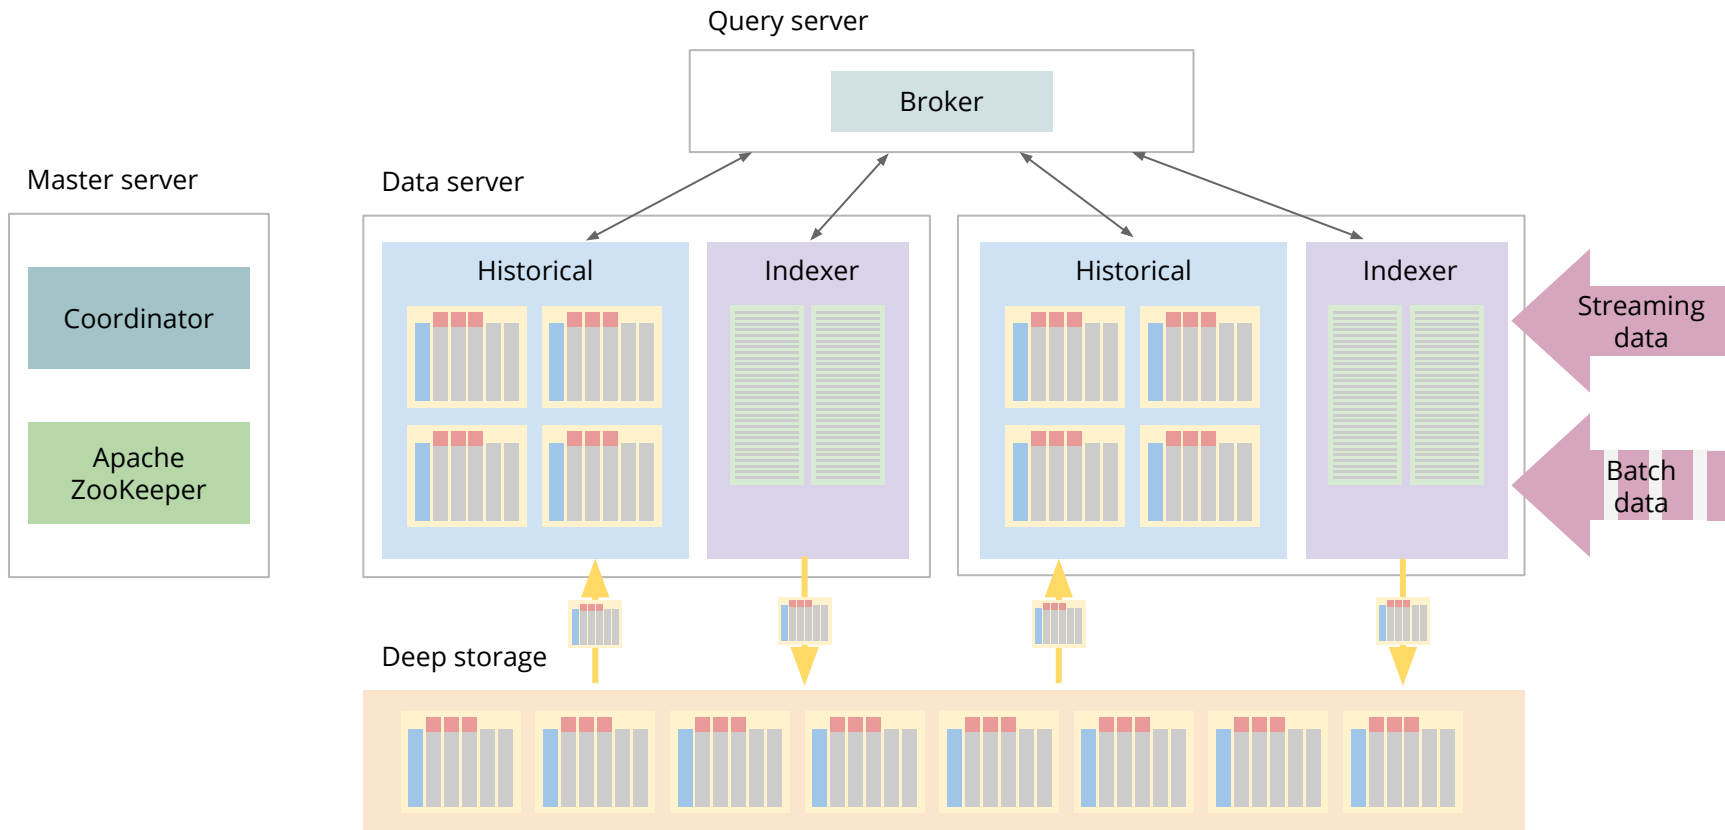


Late materialization

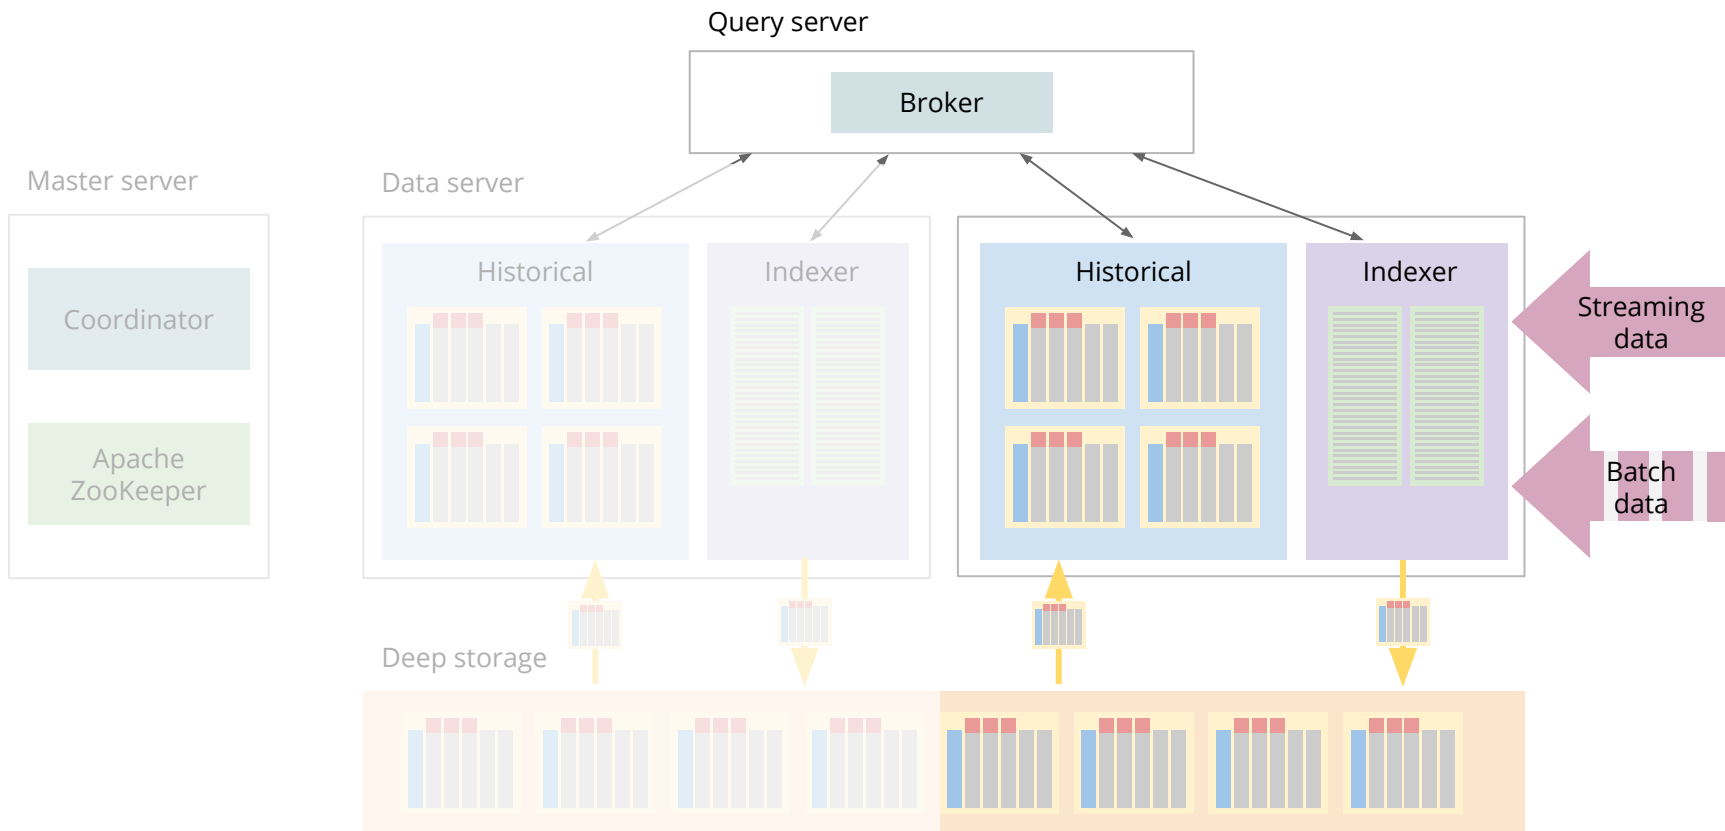
Interactive query speeds



Always online



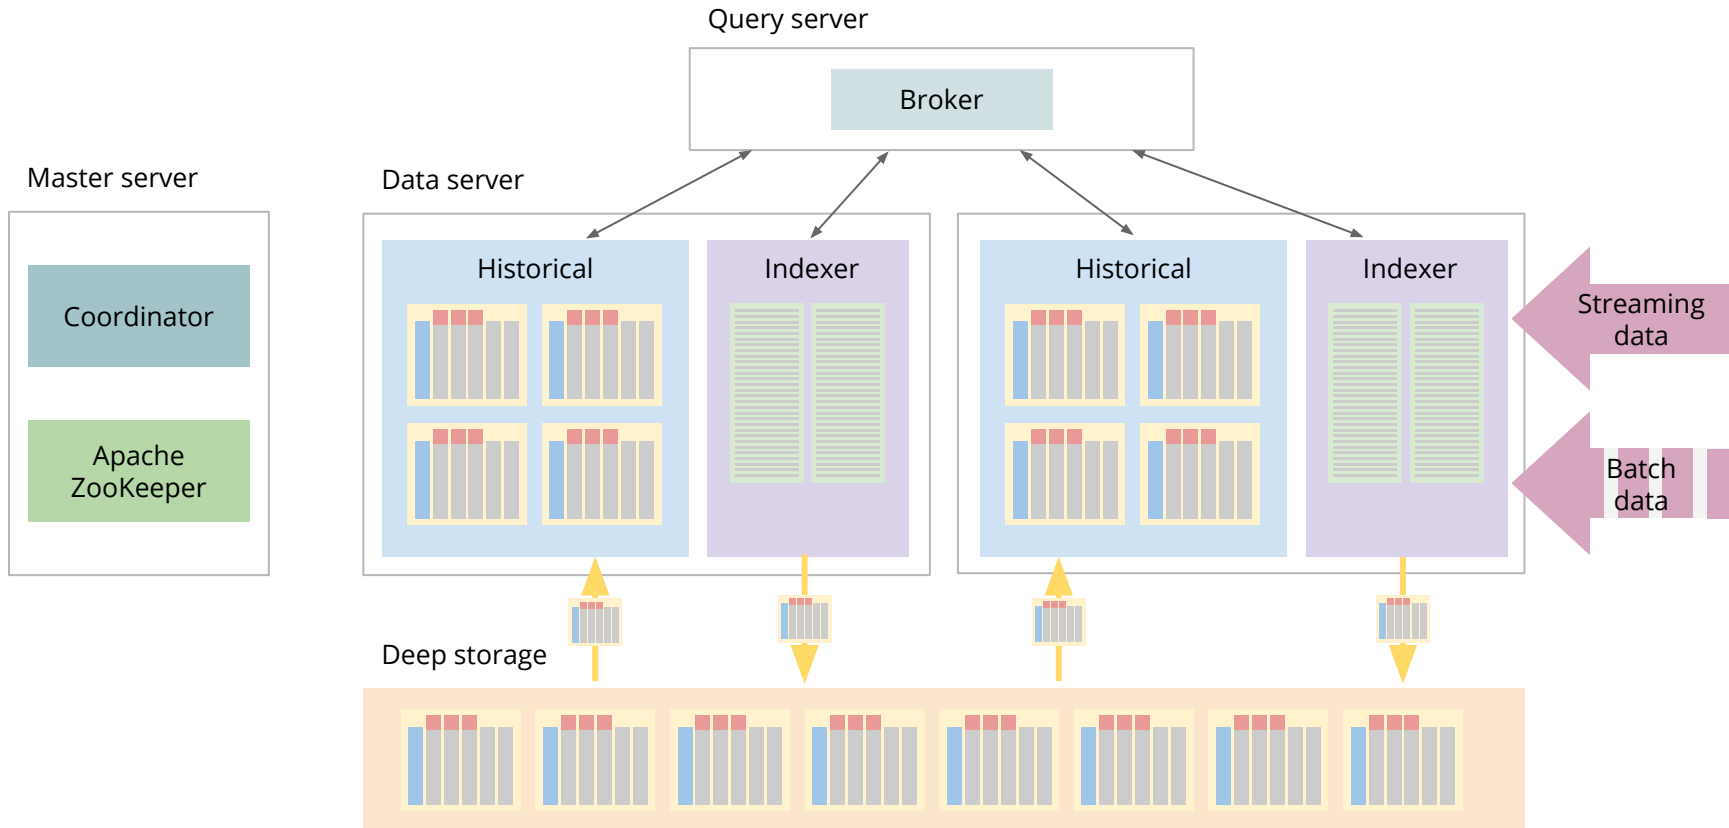
Fresh data from streams



Quality of service

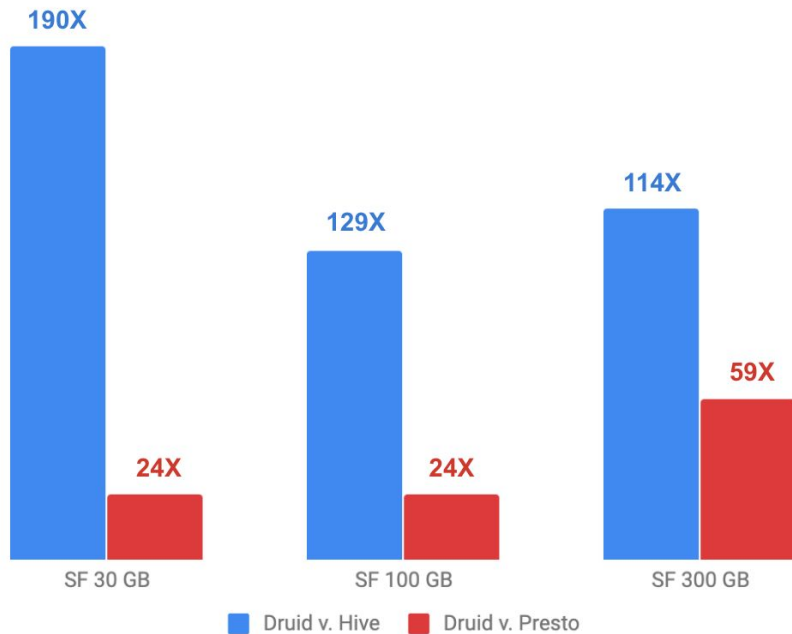


Quality of service



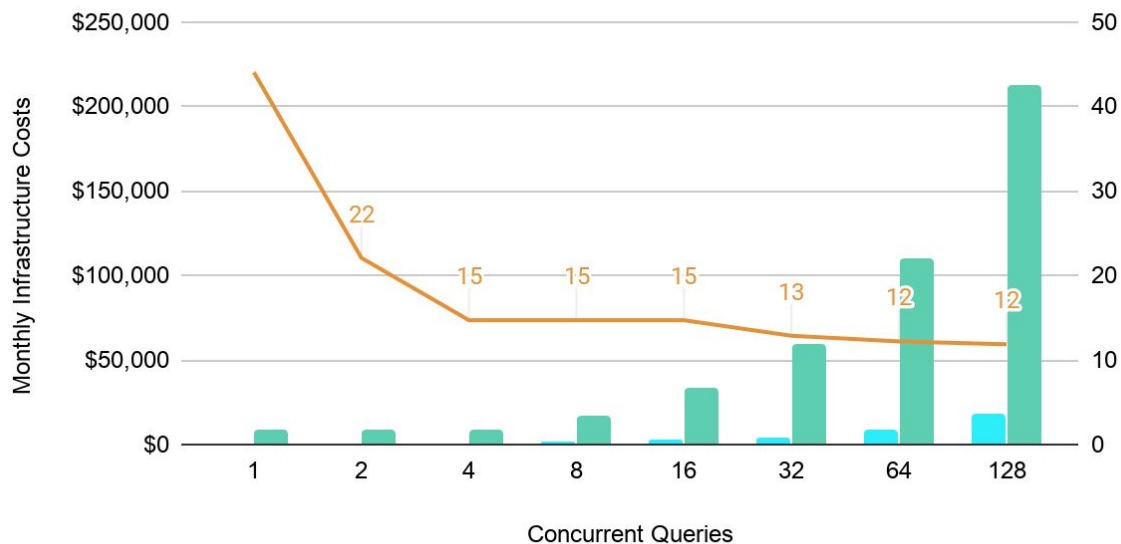
Price/performance

vs. leading open source SQL engines



Price/performance

vs. a leading cloud-based data warehouse



Wrapping up



VIRTUAL
DRUID SUMMIT

Stay in touch

Follow the Druid project on Twitter!

 @druidio

Join the community
(Mailing lists, Slack, meetups)

<https://druid.apache.org/community/>

Thank you!

Time for questions

@gianmerlino



VIRTUAL
DRUID SUMMIT